

02 / 05 DELIVERABLE · OUTPUT CONSISTENCY

# Output Consistency & Quality System

"AI slop" isn't a mystery — it's what happens when nobody's told the AI what "good" means for *this* brand. This deliverable encodes that definition and checks against it, twice, before anything ships.

Status v1 · skills ready, standards frames awaiting extraction    Owner Chii / Chiibitsu Labs

## The root cause

Generic AI output reads generic because nothing told it not to. Fix that by writing down, explicitly, what good already looks like — then have every draft checked against that definition before a human ever has to catch it by eye.

## The encoded standard

Four files, each with a different correct source and owner — deliberately pulled from where the real answer actually lives, not guessed:

| FILE                                       | ANSWERS                                       | CORRECT SOURCE                                                      |
|--------------------------------------------|-----------------------------------------------|---------------------------------------------------------------------|
| what-good-looks-like.md                    | What does a strong report actually contain?   | 2–3 gold reports                                                    |
| house-voice.md                             | How does Amplifi actually sound?              | Same gold reports                                                   |
| report-template-rules.md                   | What does the template require, structurally? | The live Canva template + brand kit — a different source on purpose |
| brand-standard.md<br><div>PER CLIENT</div> | What does this specific client expect?        | Kickoff material for that client                                    |

## Three mechanisms, one pipeline

#### DRAFT

amplifi-insights skill writes against the encoded standard — not a blank prompt.

#### QA · PASS 1

amplifi-qa checks the full markdown draft before it goes anywhere near Canva.

#### QA · PASS 2

The same gate, re-run against the visually rendered Canva export.  
**This is the only pass that actually clears a report to ship** — layout drift, stale charts, and wrong logos all happen after pass 1 already ran.

#### CAPTURE

amplifi-improve closes the loop — what this session taught goes back into the standard.

---

## Why two QA passes, not one

A markdown draft can pass every check and still ship with a stale chart or an off-brand layout once it's actually assembled in Canva — because that assembly step happens *after* the first check runs. The second pass, against the rendered deck, is the one checkpoint that's allowed to actually clear a report for the client. Anything it flags gets a human clear-or-rerun decision, every time.

---

## The loop-closer

amplifi-improve is what keeps this from calcifying. It's the same skill that runs Deliverable 1's capture loop — every session's real lessons get written down, and the good ones get promoted into these standards files weekly, so the definition of "good" keeps sharpening instead of going stale the day it was written.

---

## Status

The skills and the QA gate are built and ready. What's still pending is content, not code: Rica extracting the actual standards frames from real gold reports (roughly 90 minutes of focused work) — tracked as Phase 1's highest-priority task on the [roadmap](#).