


AI Ecosystem Architecture Map

What each tool actually does, where AI enters the workflow, and where a human always stays in charge — laid out as one map instead of six separate mental models.


Status v1 Built by Chii / Chiibitsu Labs

The four layers



The real workflow — mapped by the team, on the wall

Twelve steps, in the analysts' own words — silently mapped by each person, then walked through together in the room. Under each step, the pain your team pinned to it that day, in red. This is the map of where it hurts, and where to look when something's slow.



The pains that span the whole workflow

Some pains don't live at one step — they run across the cycle. These are the six themes your team's sticky notes grouped into, and where each one gets answered.

Brief & input quality

Answered by the client FAQ and living brief in the knowledge base → D1

Alignment gaps

Answered by one written standard everyone reads → D1 + D2

AI output quality

Answered by the encoded standard and two quality checks → D2

Data & pulling bottlenecks

Named as the first automation candidate once the data confirms it → Roadmap

Insight that doesn't compound

Answered by the insight log that accumulates per client → D1

Supplier dependency

Out of AI scope on purpose — parked as Michele's process track → Roadmap

Where AI enters, and where it doesn't

AI drafts and checks

First-pass writing, and both quality gates. Anywhere the output can be measured against an explicit, written standard.

Humans always own

Data gathering, client relationship, final visual assembly, and every ship decision. Nothing ships without a human clearing it — including what the AI-QA gate itself flags.

Tools, and what they actually do for you

TOOL	FUNCTION
Claude (Enterprise)	Runs the four skills — onboard, draft, QA, capture
Google Drive	The vault — where your knowledge base lives and everyone reaches it
Sentimo + MCP	Data connection into the drafting step
Canva FRAGILE	Final visual assembly — flagged below as a named risk
Telegram + capchecker	Daily capacity check-in — the instrument's first feed
MS Office	Supporting document work where needed

Four risks, named on purpose

Canva fragility. Crashes and access issues are real and already happening. Not replaced in this build — see the Roadmap for why, and what would change that call.

Supplier and vendor pain. Real, but a procurement problem, not an AI problem. Explicitly out of Core Build scope.

Claude usage caps. A ceiling worth watching as adoption scales past the first client.

A knowledge base that goes stale. A vault that stops getting updated is worse than no vault — this is exactly what the capture loop in Deliverable 1 exists to prevent.

Chiibitsu Labs — more human, by design.