03 / 05 DELIVERABLE · ARCHITECTURE

AI Ecosystem Architecture Map

What each tool actually does, where AI enters the workflow, and where a human always stays in charge — laid out as one map instead of six separate mental models.

Status v1 **Built by** Chii / Chiibitsu Labs

The four layers

1 · KNOWLEDGE BASE
The knowledge foundation — client briefs, standards, the whole KB. Deliverable 1.

→

2 · WORKFLOW
The 12-step path the work actually travels, in the team's own words.

→

3 · CONSISTENT OUTPUT
Lens-produced, standard-checked, on-brand. Deliverable 2.

→

4 · THE INSTRUMENT
Reads the whole system to know when it's time to hire. Deliverable 5.

The real workflow — mapped by the team, on the wall

Twelve steps, in the analysts' own words — silently mapped by each person, then walked through together in the room. Under each step, the pain your team pinned to it that day, in red. This is the map of where it hurts, and where to look when something's slow.

01
Brief
Late, incomplete briefs; the same questions asked every cycle

02
Baseline
Custom asks from the client, handled from scratch each time

03
Internal Alignment
Report drifting from the brief; no master doc to align against

04
Present
Client context scattered across past decks and heads

05
Revisions
Rounds repeat because the correction never gets written down

06
Implement
Supplier dependency — late replies, late samples, slow updates

07
Monthly Analysis
Long queue for data pulling; tool limits mid-analysis

08
Insights
Reports not evolving; no actionable recommendations

09
Report
AI slop, hallucination, templates off-brand

10
Internal Alignment
Alignment scheduled late, after the work is already done

11
Client Presentation
Manual polish so the deck "doesn't look AI"

12
Repeat
Lessons from this cycle don't carry into the next

The pains that span the whole workflow

Some pains don't live at one step — they run across the cycle. These are the six themes your team's sticky notes grouped into, and where each one gets answered.

Brief & input quality

Answered by the client FAQ and living brief in the knowledge base → D1

Alignment gaps

Answered by one written standard everyone reads → D1 + D2

AI output quality

Answered by the encoded standard and two quality checks → D2

Data & pulling bottlenecks

Named as the first automation candidate once the data confirms it → Roadmap

Insight that doesn't compound

Answered by the insight log that accumulates per client → D1

Supplier dependency

Out of AI scope on purpose — parked as Michele's process track → Roadmap

Where AI enters, and where it doesn't

AI drafts and checks

First-pass writing, and both quality gates. Anywhere the output can be measured against an explicit, written standard.

Humans always own

Data gathering, client relationship, final visual assembly, and every ship decision. Nothing ships without a human clearing it — including what the AI-QA gate itself flags.

Tools, and what they actually do for you

TOOL	FUNCTION
Claude (Enterprise)	Runs the four skills — onboard, draft, QA, capture
Google Drive	The vault — where your knowledge base lives and everyone reaches it
Sentimo + MCP	Data connection into the drafting step
Canva FRAGILE	Final visual assembly — flagged below as a named risk
Telegram + capchecker	Daily capacity check-in — the instrument's first feed
MS Office	Supporting document work where needed

Four risks, named on purpose

Canva fragility. Crashes and access issues are real and already happening. Not replaced in this build — see the Roadmap for why, and what would change that call.

Supplier and vendor pain. Real, but a procurement problem, not an AI problem. Explicitly out of Core Build scope.

Claude usage caps. A ceiling worth watching as adoption scales past the first client.

A knowledge base that goes stale. A vault that stops getting updated is worse than no vault — this is exactly what the capture loop in Deliverable 1 exists to prevent.

Chiibitsu Labs — more human, by design.